

# What you don't know about containers *(and why it matters)*

DOAG CloudLand 2024  
Brühl, Germany  
June 20, 2024



**Sean Scott**



# Database Reliability Engineering

## Business Continuity :: HA & DR

## Automation :: Observability

Real Application Clusters :: Data Guard :: Sharding

Containerization :: Terraform :: Ansible

Exadata & Engineered Systems

AHF :: TFA :: GIMR :: CHA

**Data on Kubernetes Community Ambassador**

**Managing Principal Consultant  
Viscosity North America**





## 500+ technical experts helping peers globally

The **Oracle ACE Program** recognizes and rewards community members for their technical and community contributions to the Oracle community



### 3 membership tiers



Oracle ACE  
Director

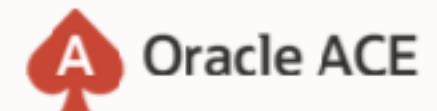


Oracle ACE  
Pro



Oracle ACE  
Associate

For more details on Oracle ACE Program:  
[ace.oracle.com](http://ace.oracle.com)



**Nominate**  
yourself or someone you know:  
[ace.oracle.com/nominate](http://ace.oracle.com/nominate)



# Oracle on Docker

Running Oracle Databases in Linux Containers

# Oracle on Docker

Running Oracle Databases in  
Linux Containers

—  
Sean Scott

apress®

Free sample chapter:  
<https://oraclesean.com>

# Lab Preparation



## You will need:

- A Linux environment
- sudo permission (to mount/umount)
- Three (or more) terminal windows
- Optional: Docker/Podman

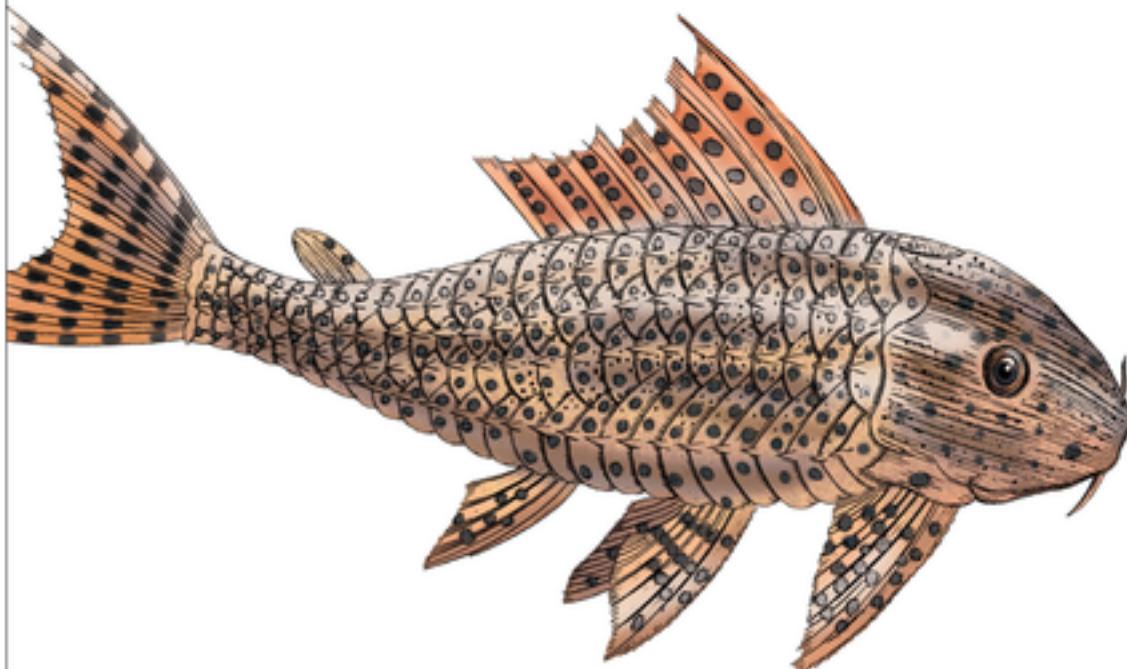
Copy/pull code examples and slides:

<https://github.com/oraclesean/cloudland24>

O'REILLY®

# Container Security

Fundamental Technology Concepts That  
Protect Containerized Applications



Liz Rice

# Container Security

Liz Rice

# What Are Containers?

Containers are like lightweight Virtual Machines, right?



# Electricity is like a water hose

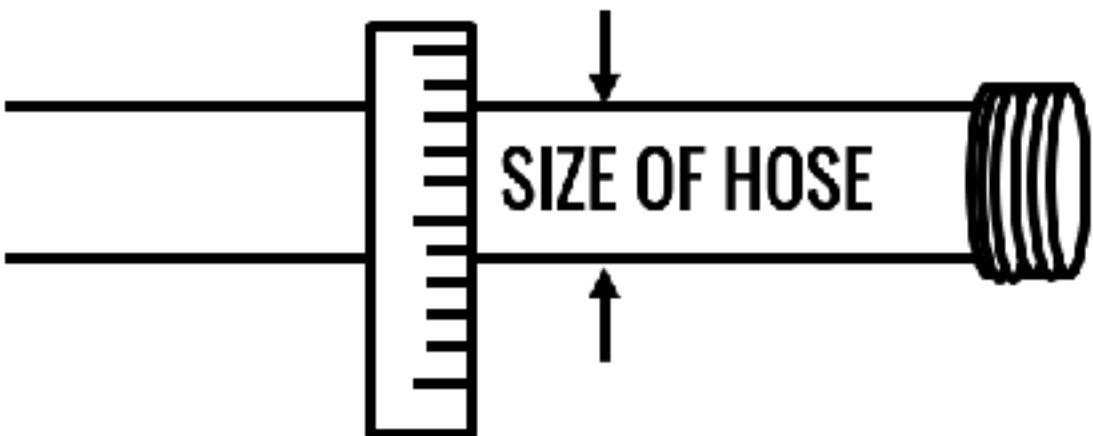
**Voltage**

Volts (V)



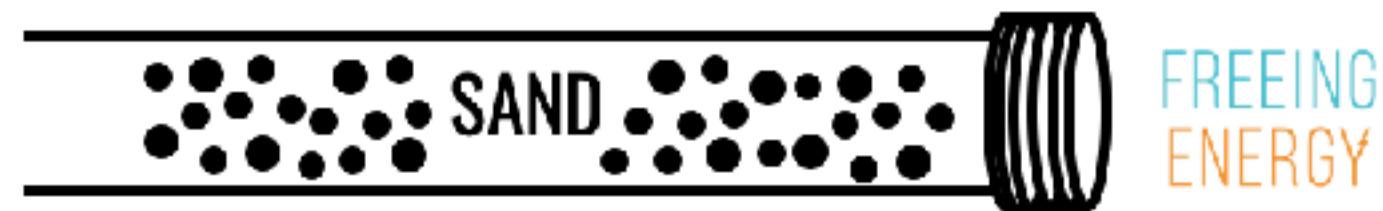
**Current**

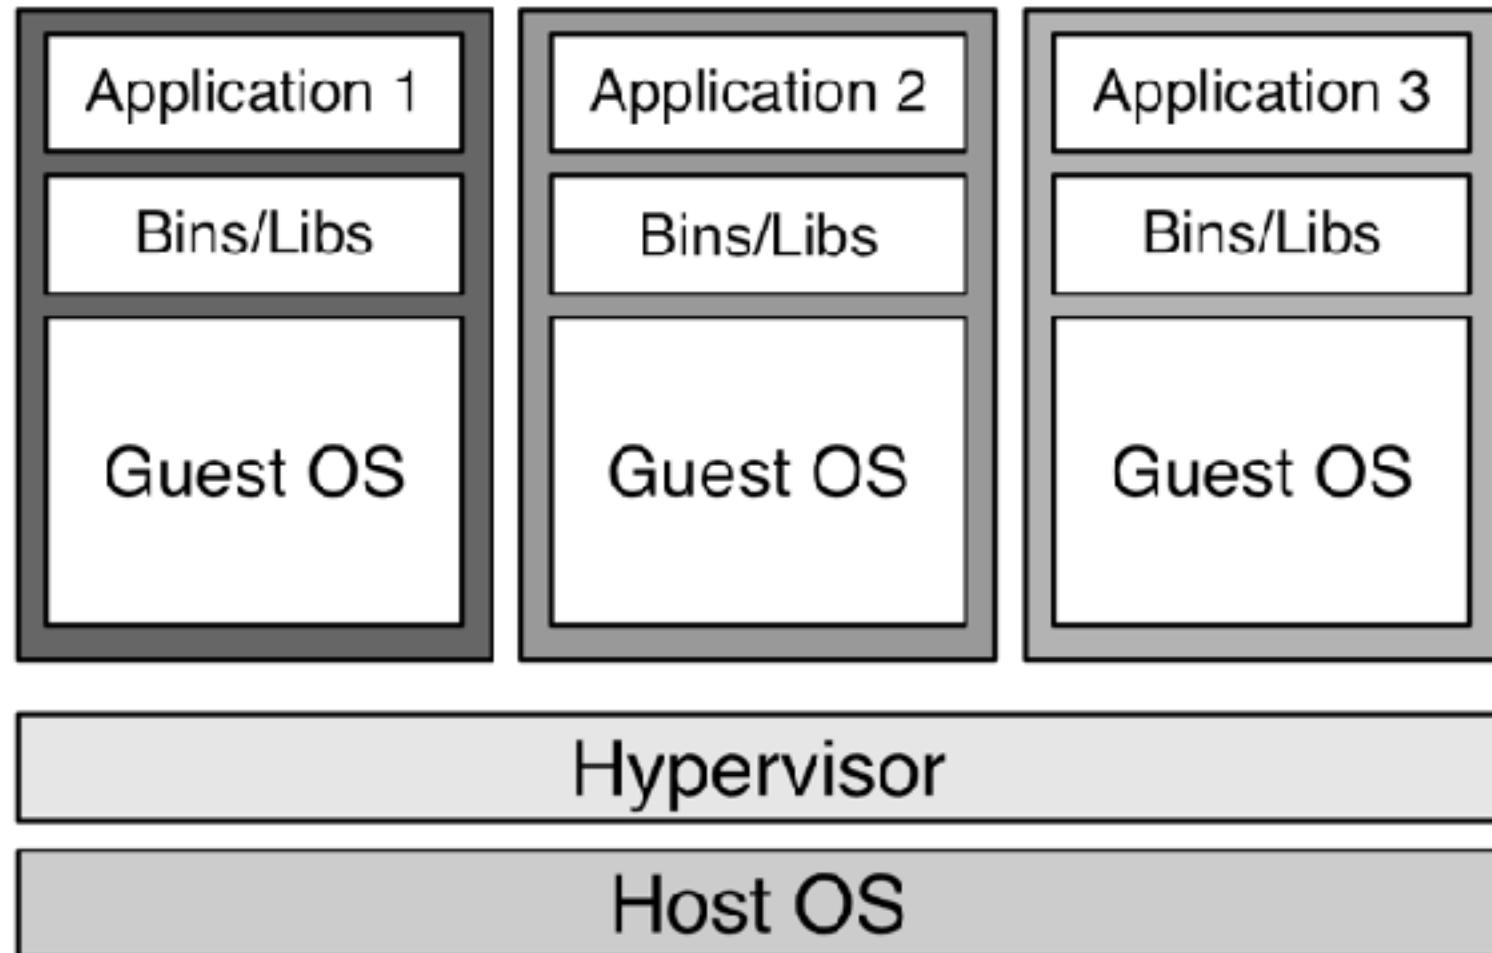
Amps (A or I)



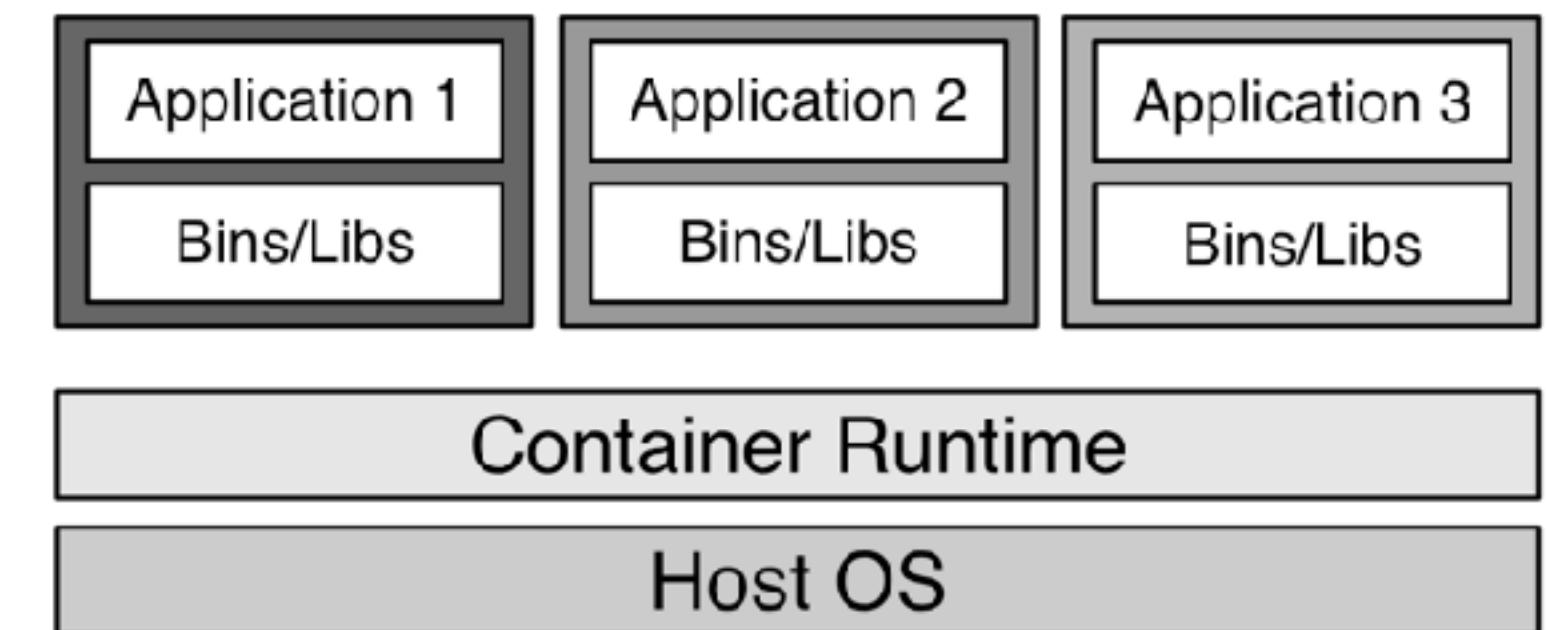
**Resistance**

Ohms (R or  $\Omega$ )





Three full operating systems



Three application/executable directories

# 13 Oracle Databases on a Laptop

#	dps	NAMES	IMAGE	PORTS	NETWORKS	STATE
		ORCL112	oraclesean/db:11.2.0.4-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL121	oraclesean/db:12.1-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL122	oraclesean/db:12.2-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL1914	oraclesean/db:19.14-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL19151	oraclesean/db:19.15.1-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL19152	oraclesean/db:19.15.2-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL1915	oraclesean/db:19.15-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL19161	oraclesean/db:19.16.1-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL1916	oraclesean/db:19.16-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL1917	oraclesean/db:19.17-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL216	oraclesean/db:21.6-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL217	oraclesean/db:21.7-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running
		ORCL218	oraclesean/db:21.8-EE	1521/tcp, 5500/tcp, 8080/tcp	oracle-db	running

2018 MacBook Pro 15", 2.2GHz 6-core Intel Core i7, 16GB, 1TB SSD

**Virtual Machines:** Bootable OS, often heavy

**Containers:** Support one application or service

**Virtual Machines:** Bootable OS, often heavy

A 100-page book of games & puzzles

**Containers:** Support one application or service

**Virtual Machines:** Bootable OS, often heavy

A 100-page book of games & puzzles

**Containers:** Support one application or service

A sheet of paper with a Tic-Tac-Toe grid

**Games** have structured playing  
surfaces and rules.

**Services** deliver games to players.

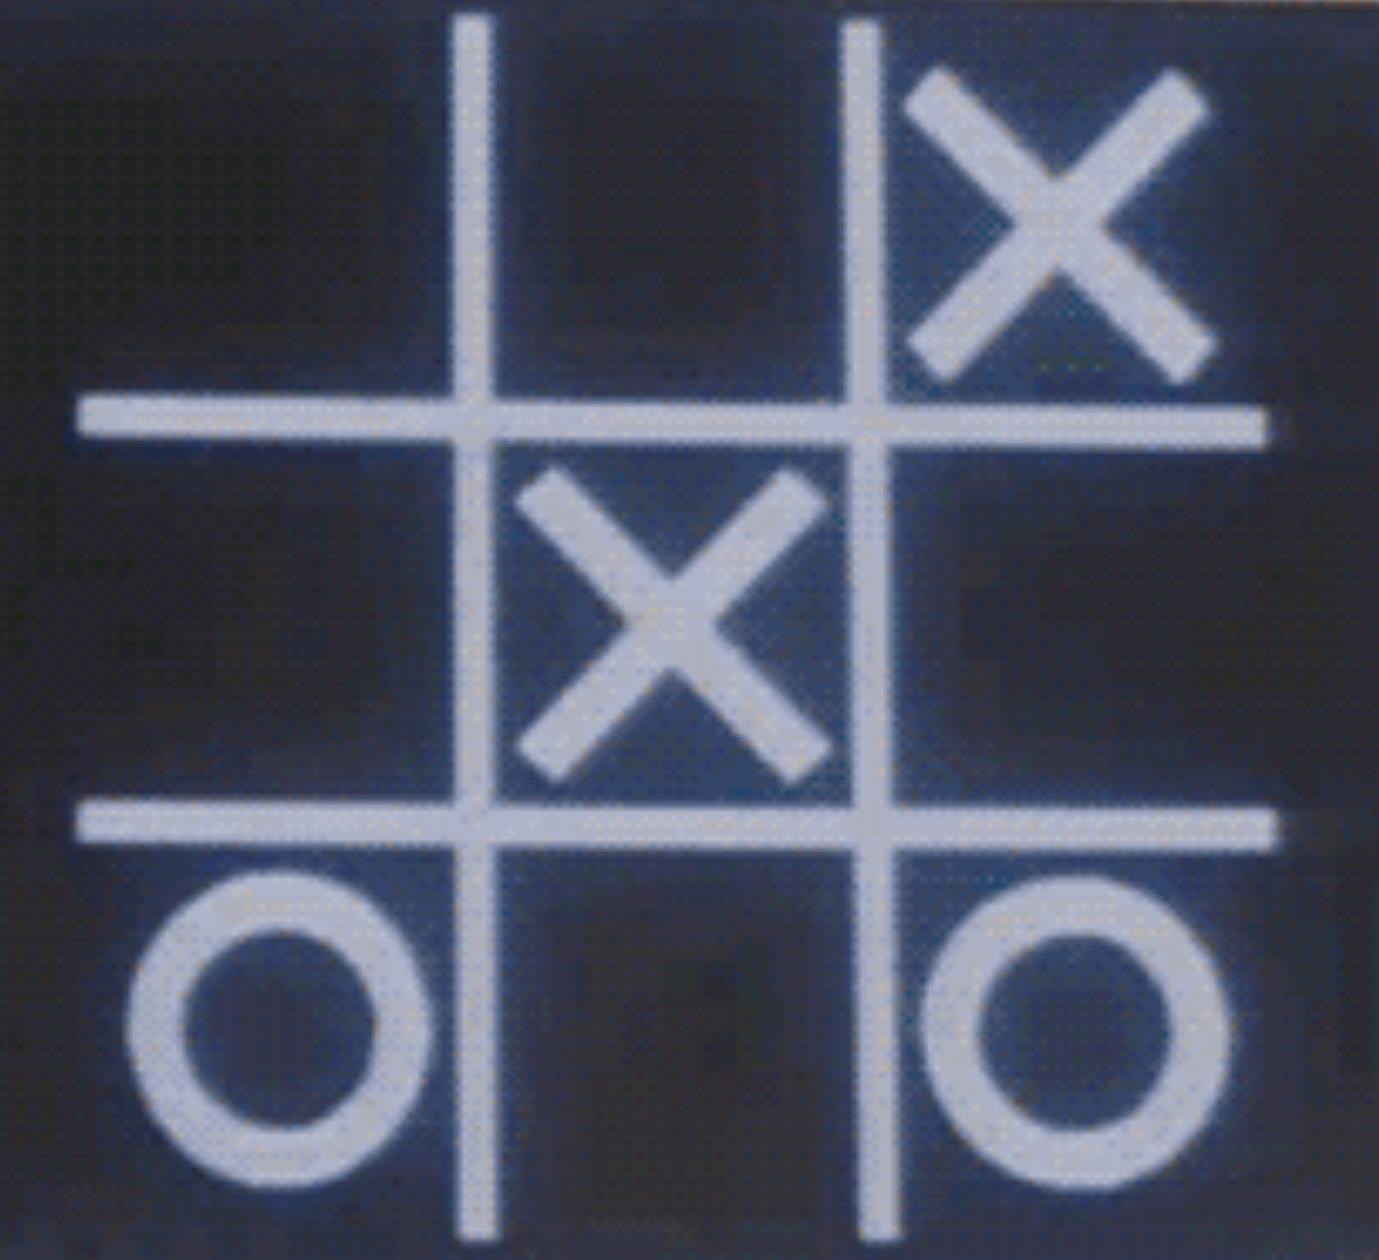
**Images** include the  
game rules and playing surface.

Images are a minimal filesystem & metadata  
for running a service.

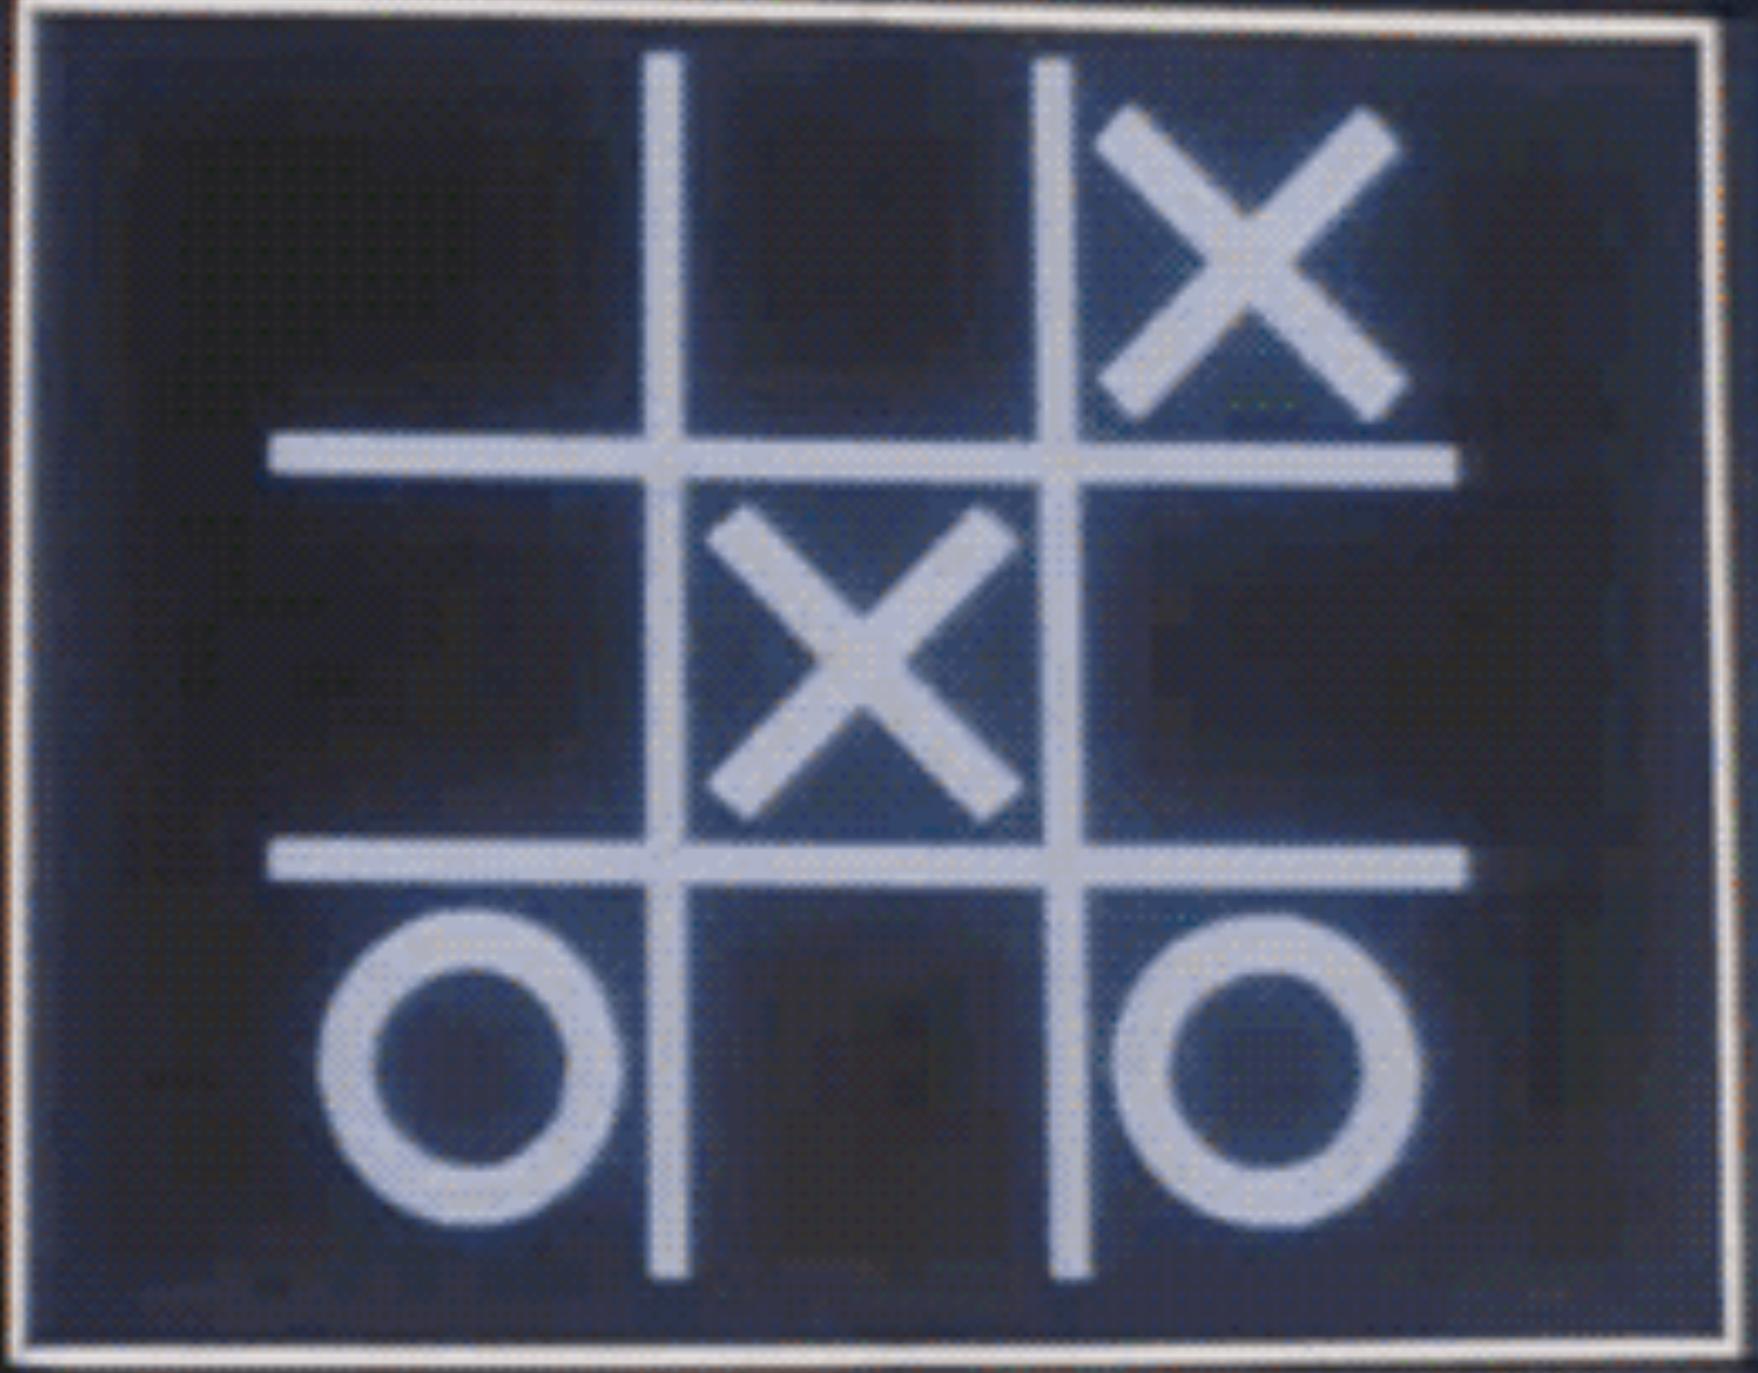
Running an image starts a **Container**.

A container is:  
**an isolated host process and  
a Union Filesystem.**

SHALL WE PLAY A GAME?



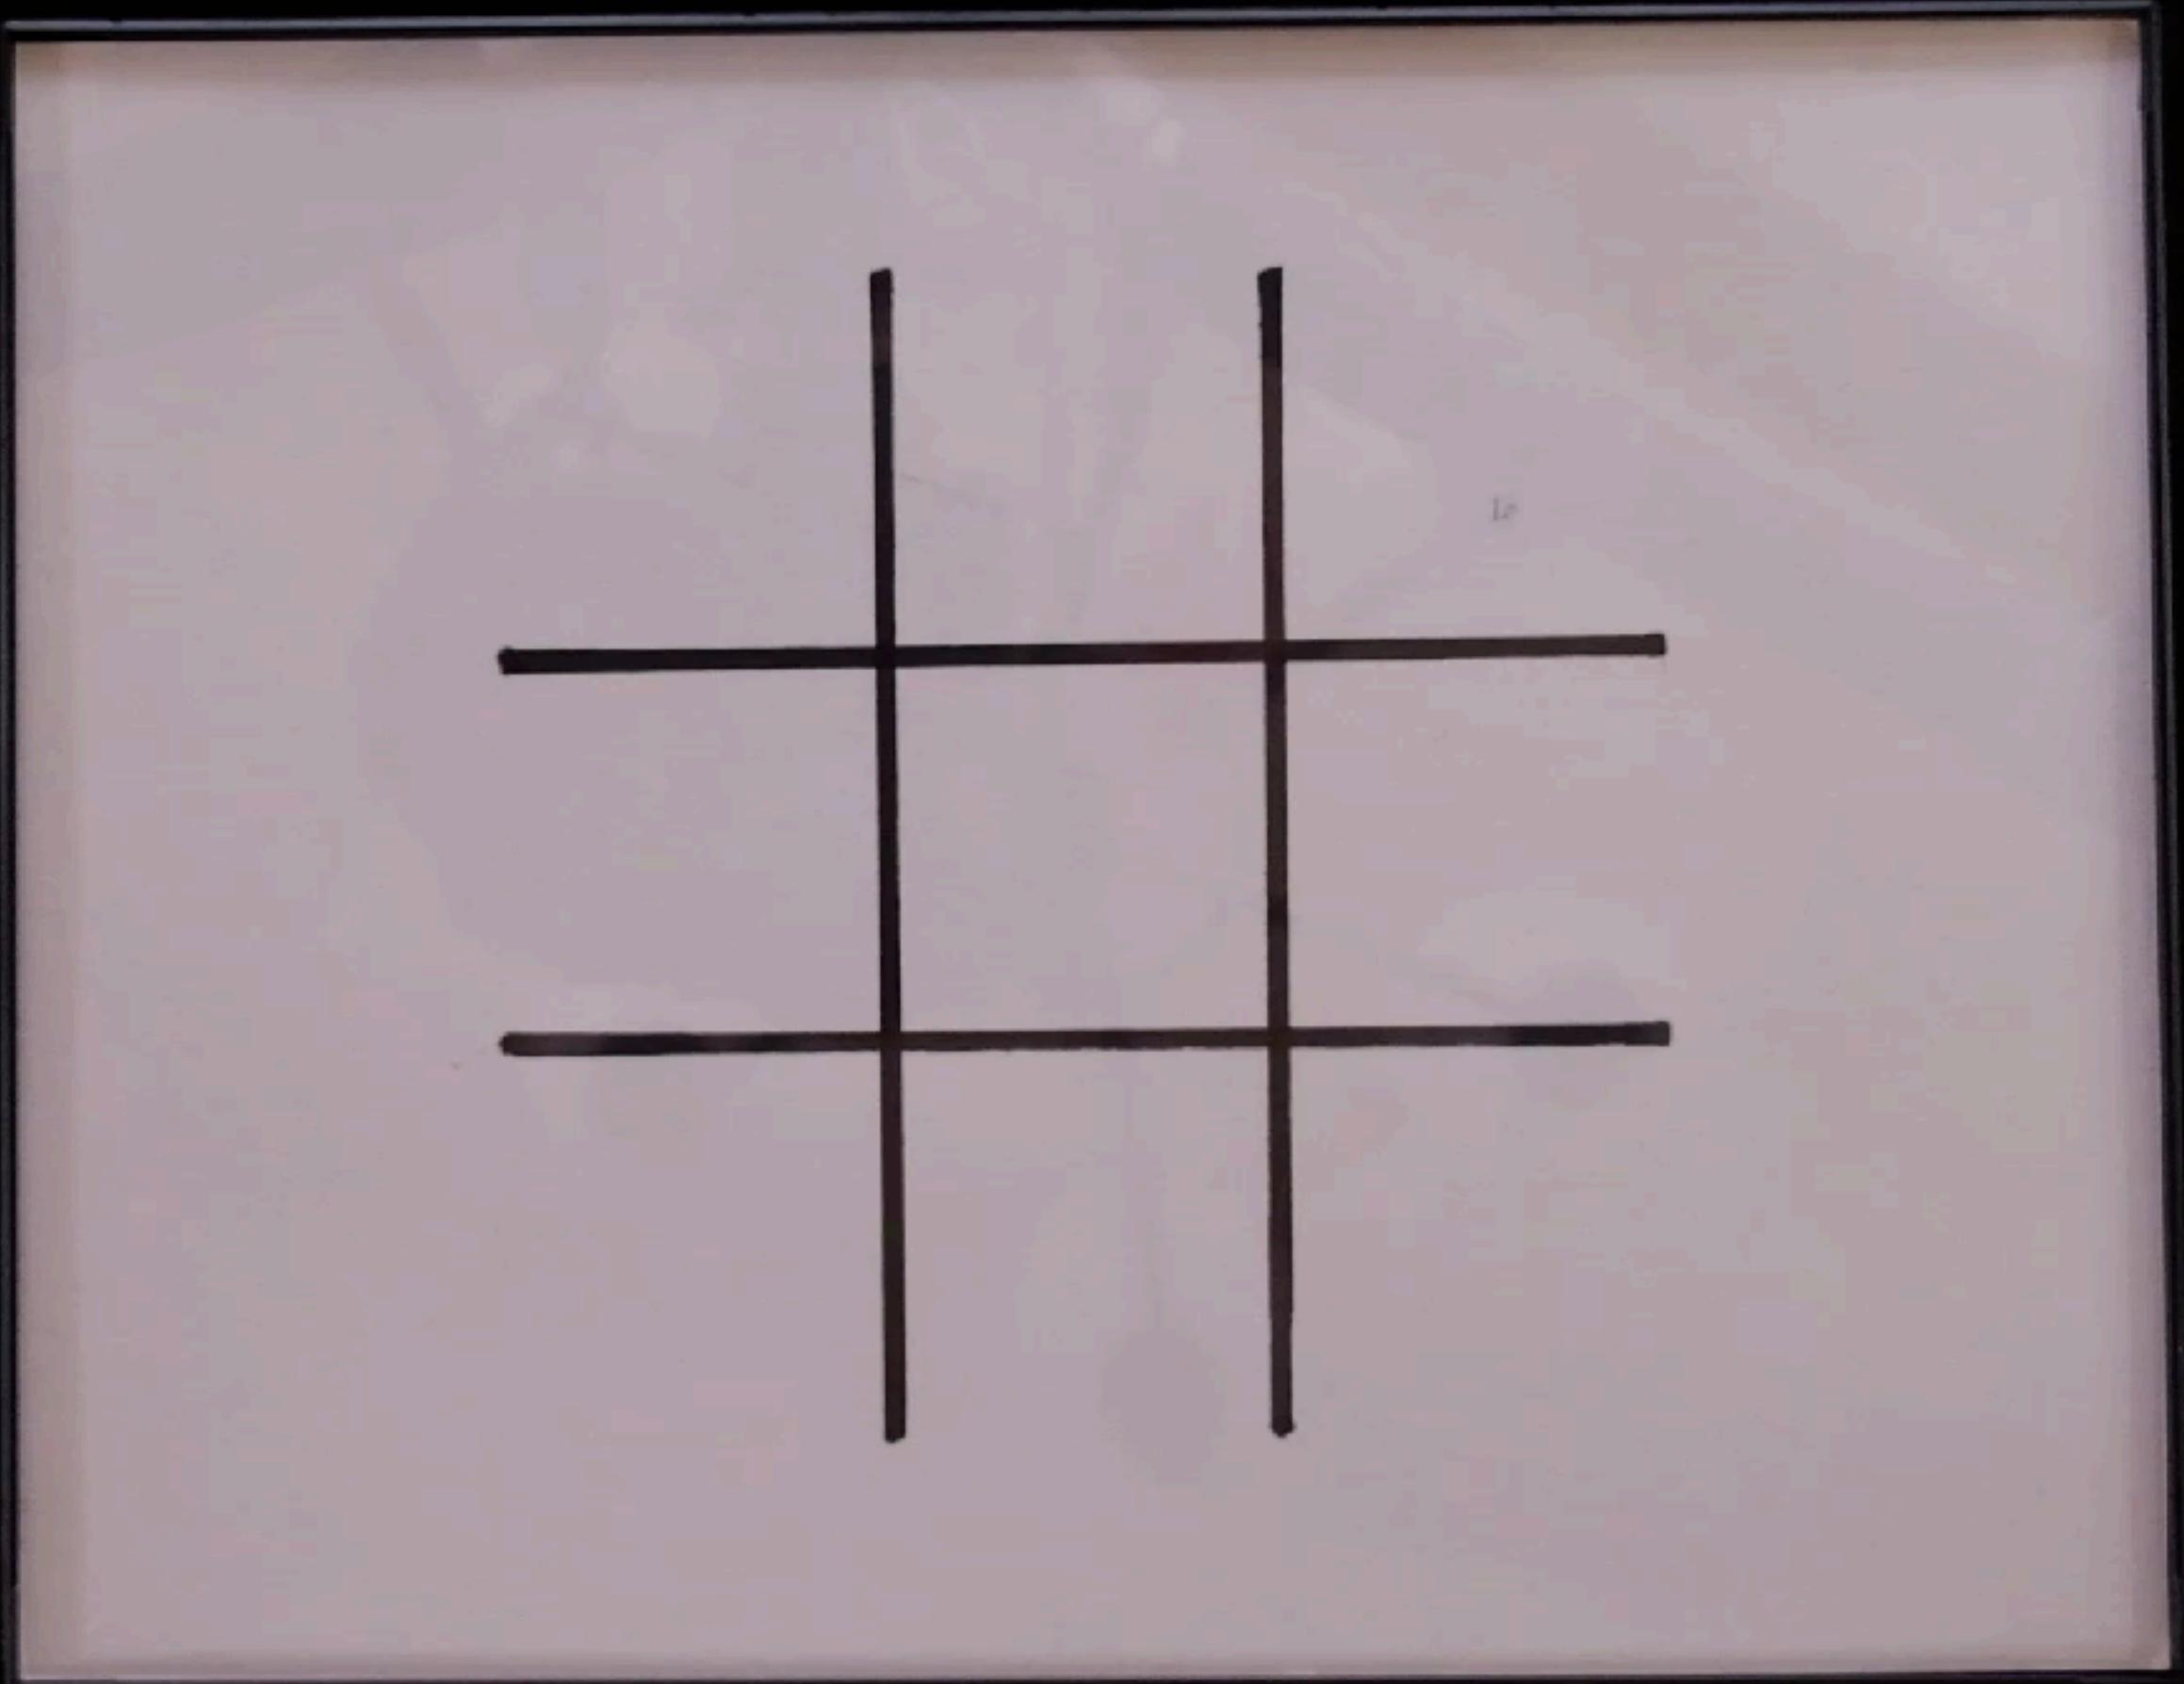
UPGRADE  
YOUR POSITION  
**9377**



UPGRADE  
YOUR POSITION  
**3652**      **9080**      **4512**      **8884**      **6668**



UPGRADE  
YOUR POSITION  
**6668**



A **Union Filesystem** has three layers:

A Merge or Union Layer.

An Upper Layer.

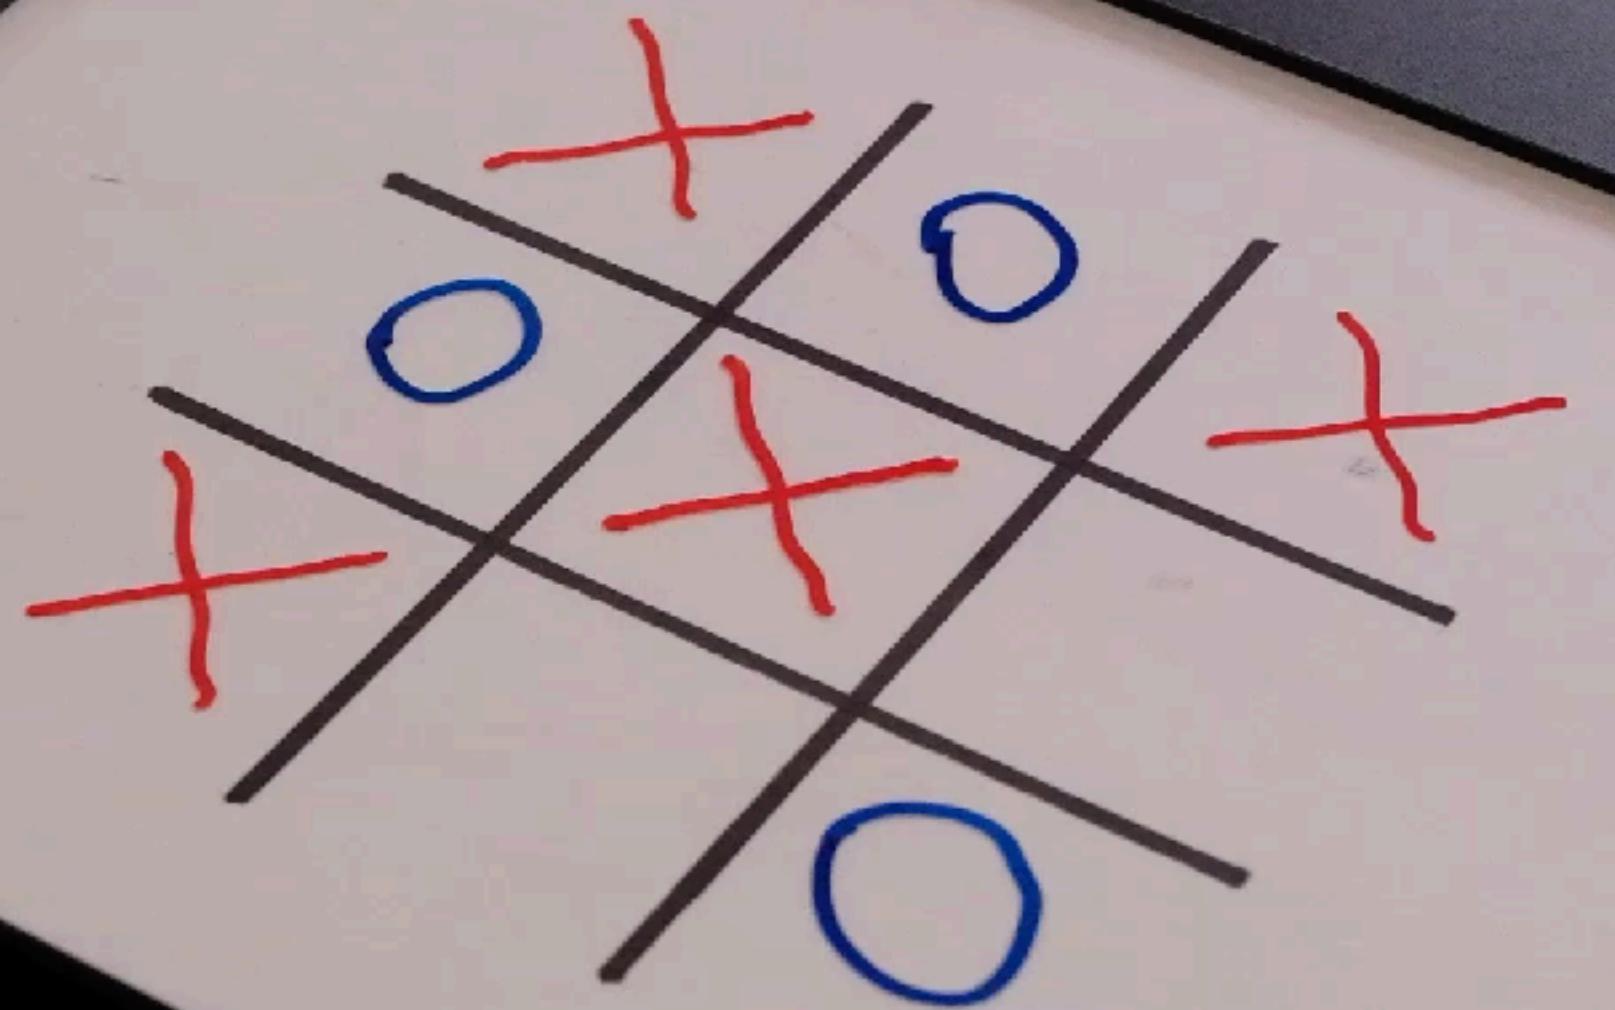
A Lower Layer.

A **Union Filesystem** has three layers:

A Merge or Union Layer: The **game**.

An Upper Layer: The **player's moves**.

A Lower Layer: The **board and rules**.



# The **Upper** (merge) layer isn't "real."

**The Upper** (merge) layer isn't "real."

It's a perceived (calculated) view  
of two overlapping layers.

The **Upper** (merge) layer isn't "real."

It's a perceived (calculated) view  
of **moves** and **the board**.

# Union Filesystem Lab



```
# Create some directories:
```

```
mkdir -p ~/workshop/  
{image,con1_upper,con1_work,container1}
```

```
# cd into the main directory:
```

```
cd ~/workshop
```

```
# Set the architecture:  
  
arch="$(arch)"  
arch="${arch/arm64/aarch64}"  
arch="${arch/amd64/x86_64}"  
  
# wget Alpine:  
  
wget -qO- https://dl-cdn.alpinelinux.org/alpine/v3.19/releases/"${arch}"/  
alpine-minirootfs-3.19.1-"${arch}".tar.gz \  
| tar xvz -C ~/workshop/image
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=con1_upper, \
workdir=con1_work \
overlay container1
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=con1_upper, \
workdir=con1_work \
overlay container1
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=con1_upper, \
workdir=con1_work \
overlay container1
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=con1_upper, \
workdir=con1_work \
overlay container1
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \  
-o lowerdir=image, \  
upperdir=con1_upper, \  
workdir=con1_work \  
overlay container1
```

```
# List files in different layers:
```

```
ls -l ~/workshop/image
```

```
ls -l ~/workshop/con1_upper
```

```
ls -l ~/workshop/container1
```

```
# Create files in different layers:
```

```
echo "AAA" > ~/workshop/image/AAA
```

```
echo "BBB" > ~/workshop/con1_upper/BBB
```

```
# Change a file in the container layer:  
echo "aaa" > ~/workshop/con1_upper/AAA
```

```
# Check the contents of file AAA in each layer:
```

```
cat ~/workshop/image/AAA
```

```
cat ~/workshop/con1_upper/AAA
```

```
cat ~/workshop/container1/AAA
```

```
# Update PATH
```

```
PATH=$PATH:/bin
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
--user --ipc --net \  
--map-root-user \  
chroot ~/workshop/container1 \  
/bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
--user --ipc --net \  
--map-root-user \  
chroot ~/workshop/container1 \  
/bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container1 \  
    /bin/sh
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container1 \  
    /bin/sh
```

# Check the hostname:

hostname

```
# Change the hostname
```

```
hostname alpine  
hostname
```

```
# Explore the "container"
```

```
cat /etc/os-release
```

```
uname -a
```

```
whoami
```

```
echo $$
```

```
pwd
```

# What files are here?

ls -l /

cat AAA

cat BBB

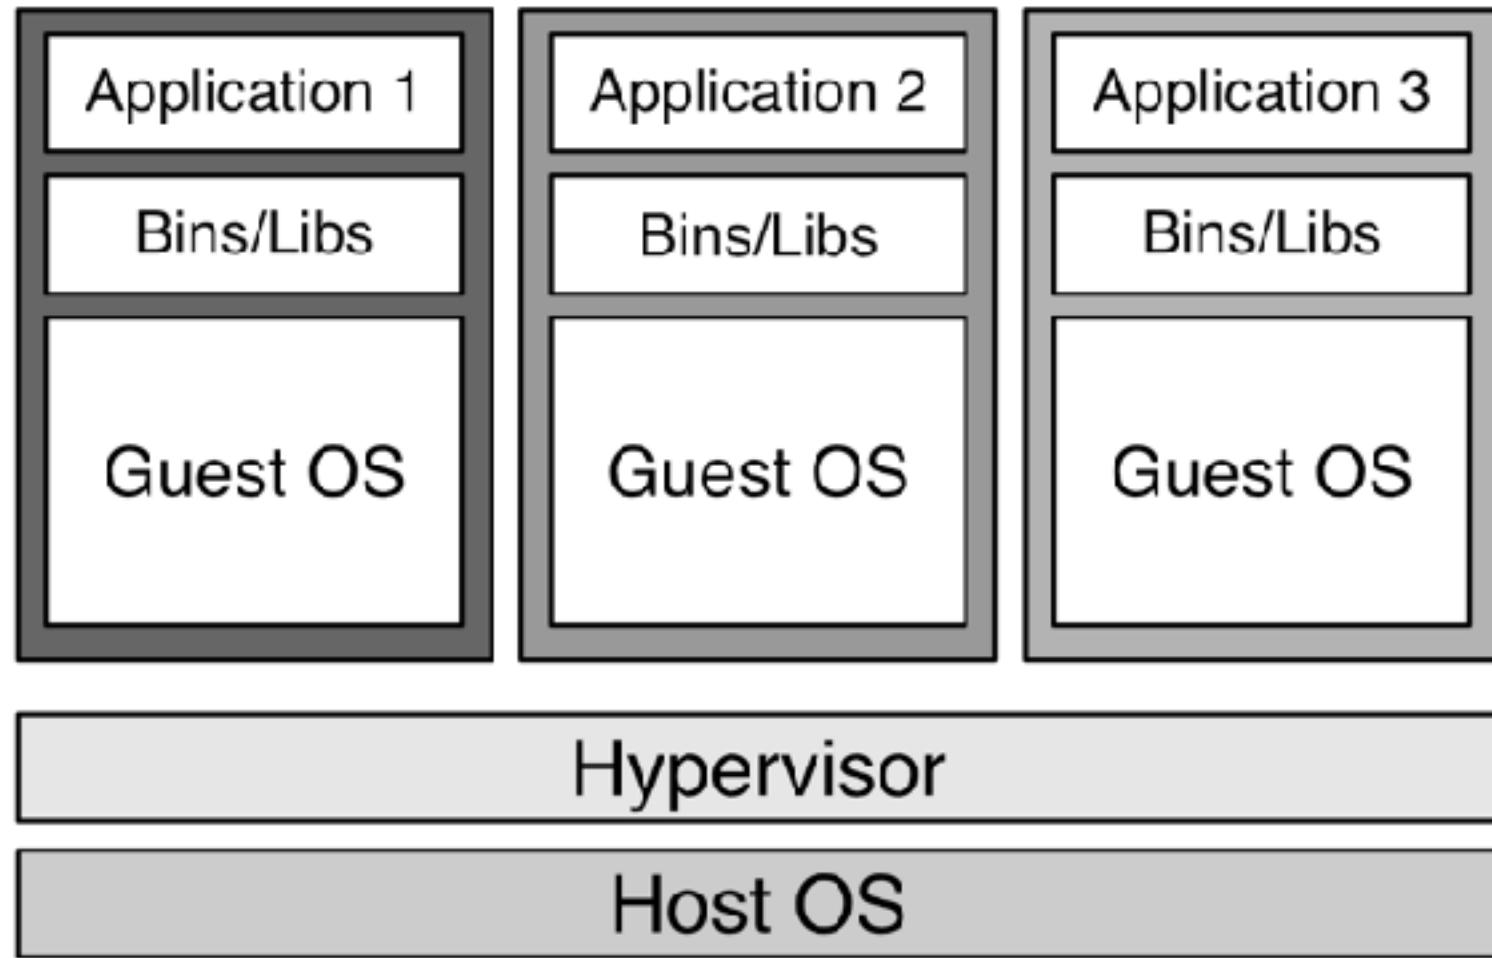
```
# Add a group and user:
```

```
addgroup -g 12345 workshop_g
```

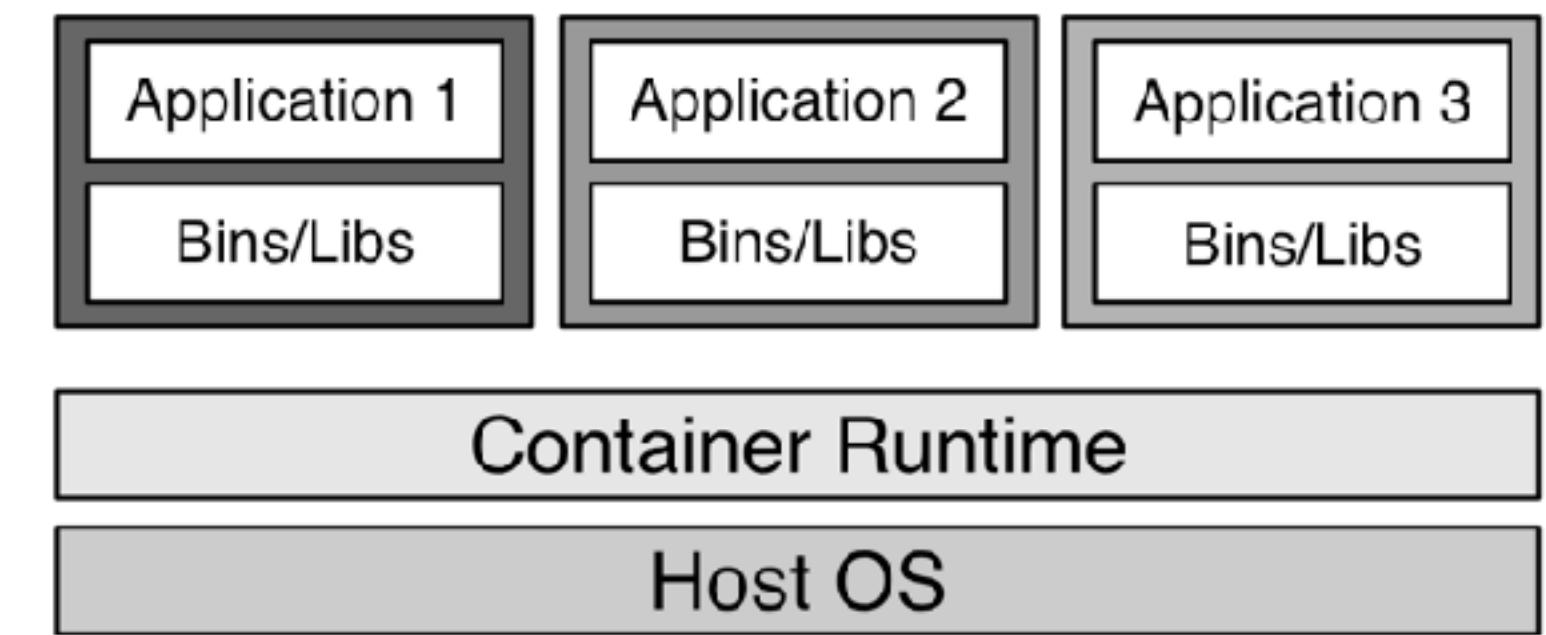
```
adduser -H -S -u 54321 workshop_u -G workshop_g
```

# Containers are Magic!

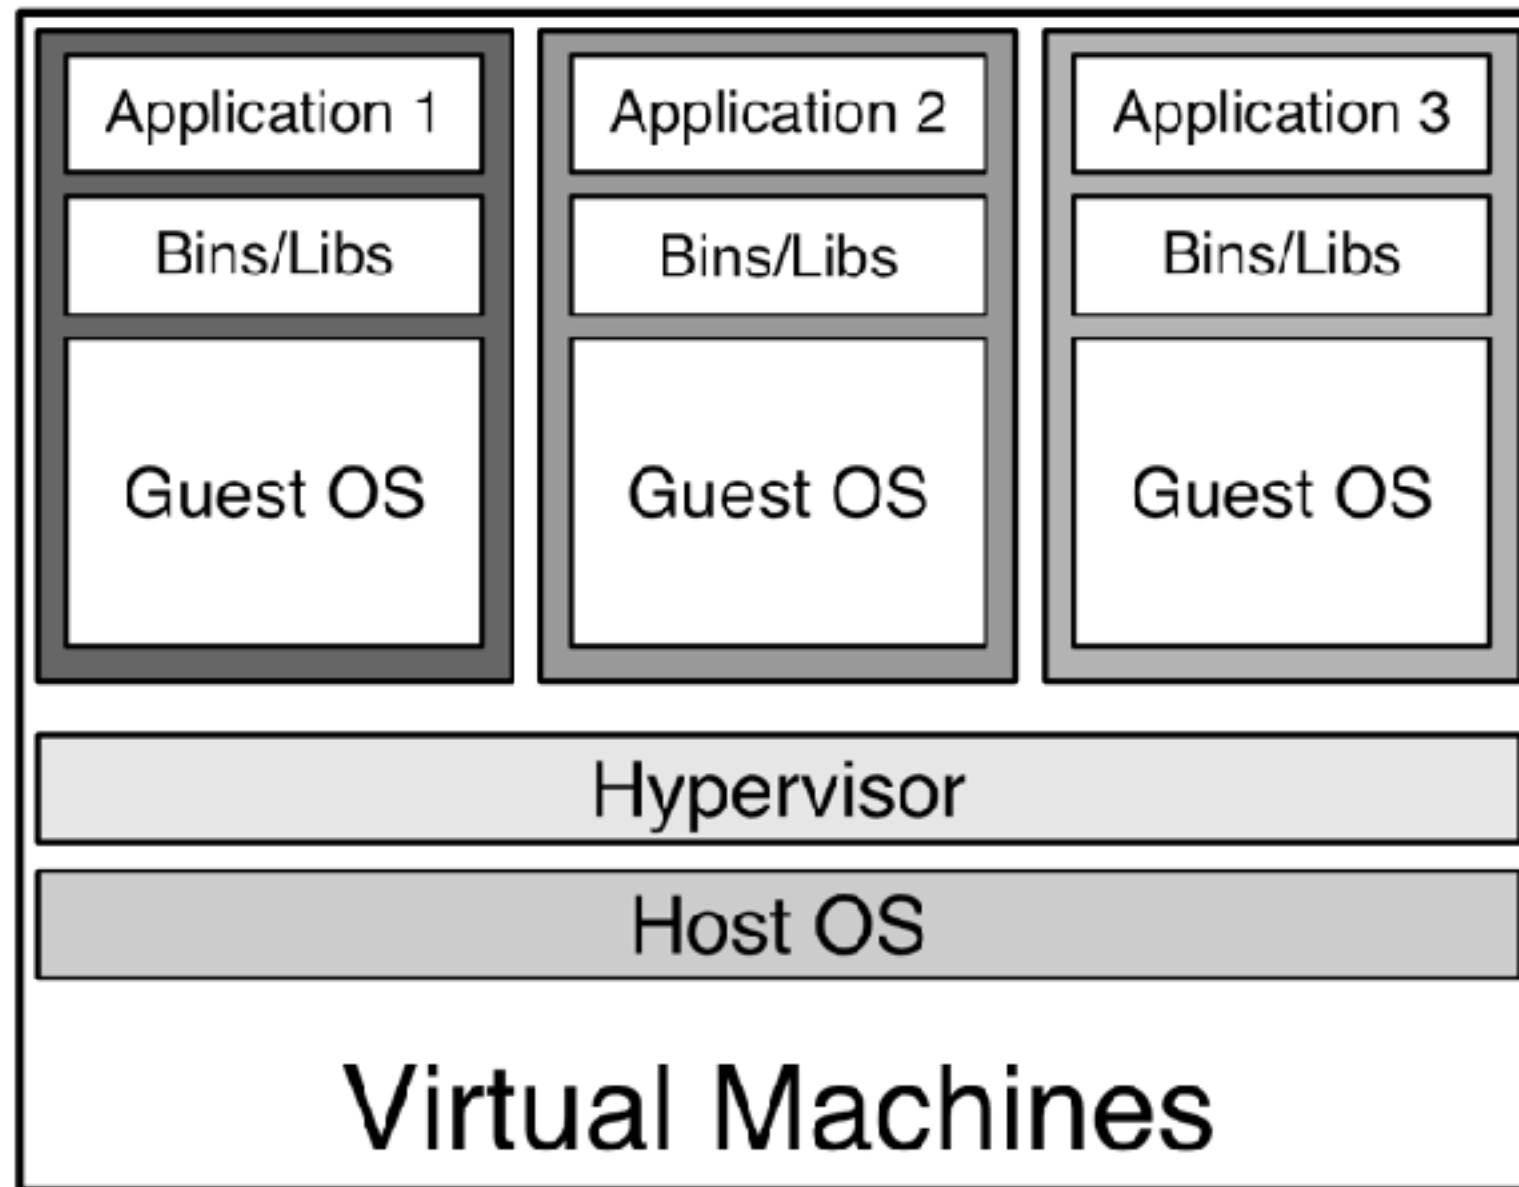




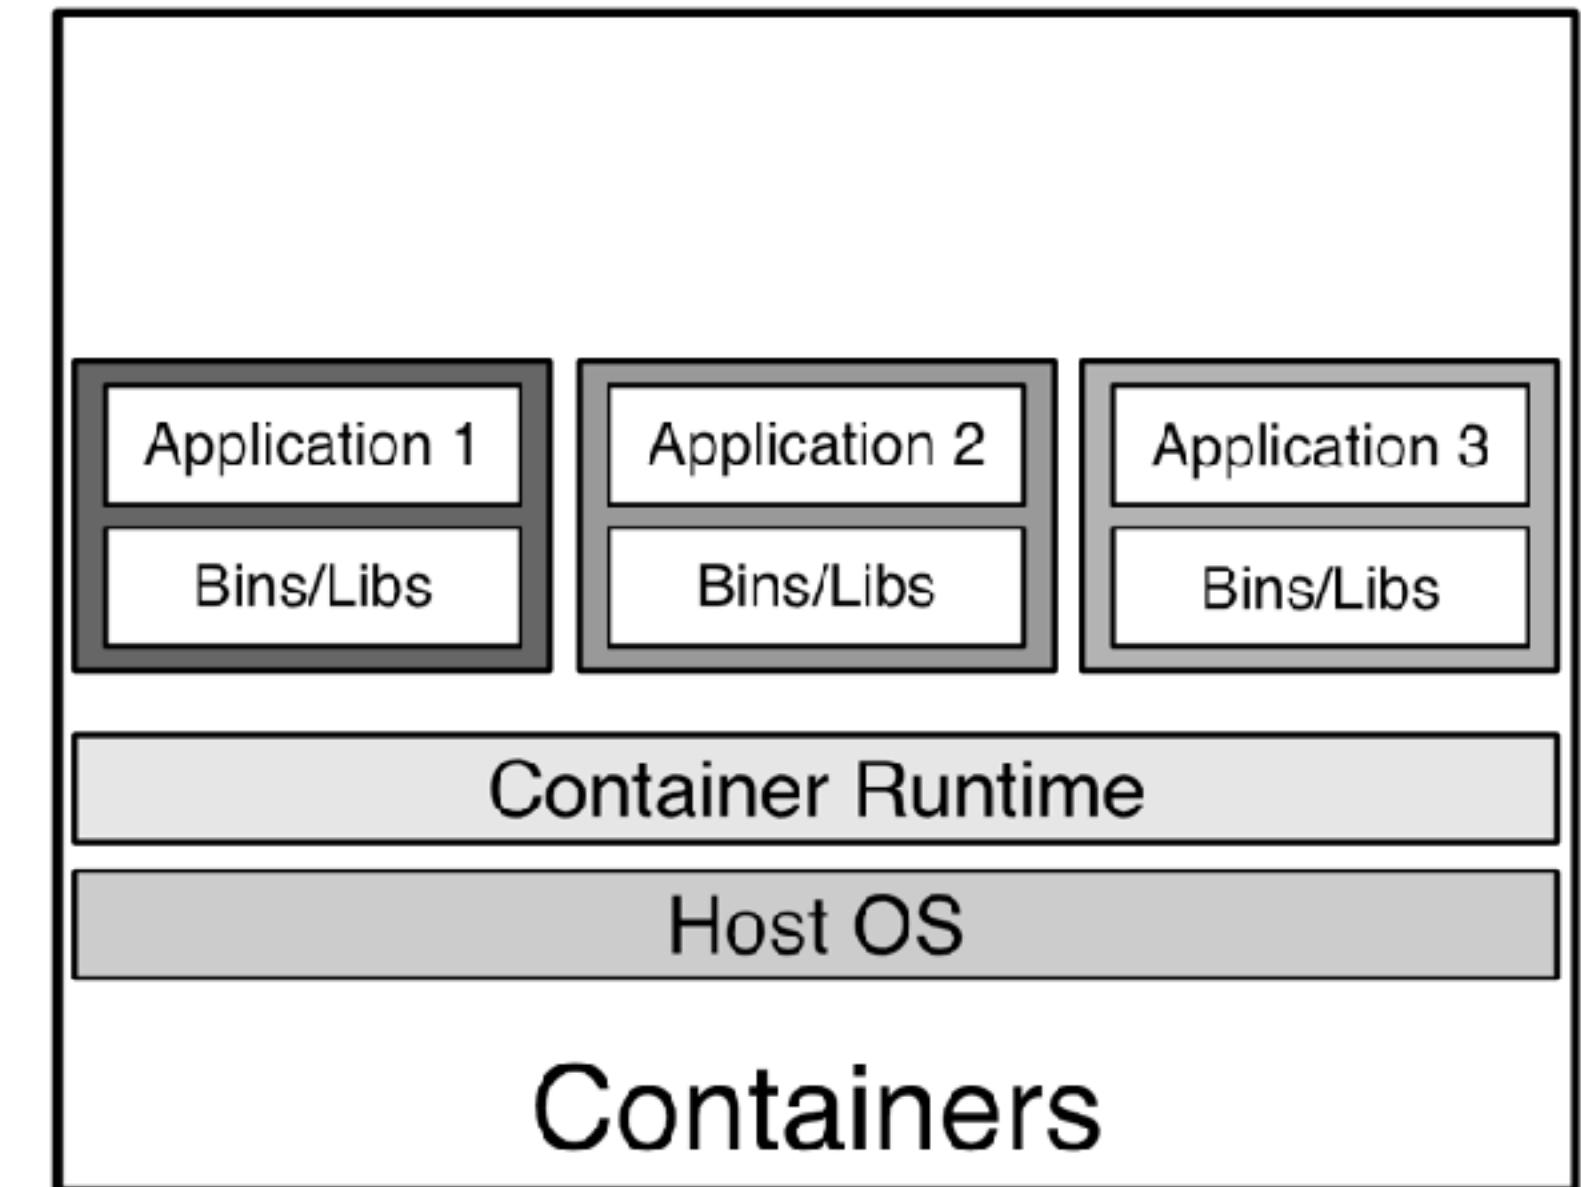
Three full operating systems



Three application/executable directories



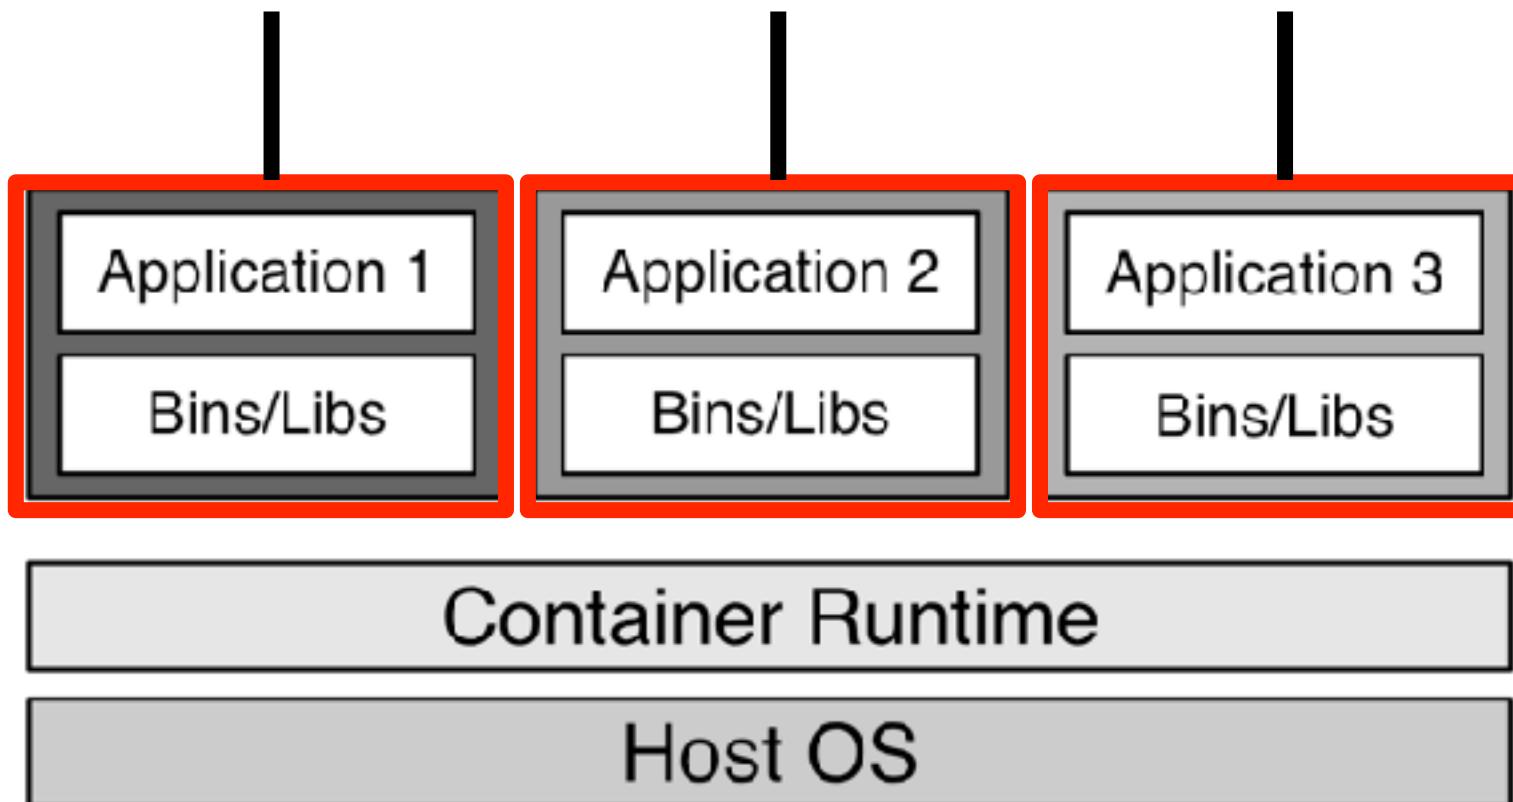
Three full operating systems



Three application/executable directories

# How much space do three containers use?

500MB 500MB 500MB



$$\begin{array}{r} 500\text{MB} \\ + 500\text{MB} \\ + 500\text{MB} \\ \hline 1,500\text{MB} \end{array}$$

Multiple identical containers will share an **image**.

```
# Prepare directories for a second container  
mkdir -p ~/workshop/{con2_upper,con2_work,container2}  
  
cd ~/workshop
```

```
# Create a second overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=con2_upper, \
workdir=con2_work \
overlay container2
```

```
# Create a second overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=con2_upper, \
workdir=con2_work \
overlay container2
```

```
# Create a second overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=image, \
upperdir=con2_upper, \
workdir=con2_work \
overlay container2
```

```
# Set the environment  
  
PATH=$PATH:/bin  
  
# "Run" the "container"  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container2 \  
    /bin/sh
```

# What files are here?

ls -l /

cat AAA

cat BBB

```
# Manipulate files within the container:
```

```
echo "111" > /AAA  
echo "222" > /BBB
```

```
# Check the contents of files from the host:
```

```
cat ~/workshop/image/AAA
```

```
cat ~/workshop/con2_upper/AAA
```

```
cat ~/workshop/con2_upper/BBB
```

Containers **persist** state (moves) in the run-time layer of its Union Filesystem.

Containers are **stateful**.

```
# Exit and restart the container:
```

```
exit
```

```
unshare --fork --uts --pid --mount \
--user --ipc --net \
--map-root-user \
chroot ~/workshop/container2 \
/bin/sh
```

```
# Confirm the file contents:
```

```
ls -l /  
cat AAA  
cat BBB
```

# Container Properties

Dropping a **Container** removes its **Layers**.



```
# "Drop" the container:
```

```
sudo umount container2
```

```
rm -fr ~/workshop/container2
```

```
rm -fr ~/workshop/con2_upper
```

```
rm -fr ~/workshop/con2_work
```

# **Containers are Ephemeral.**

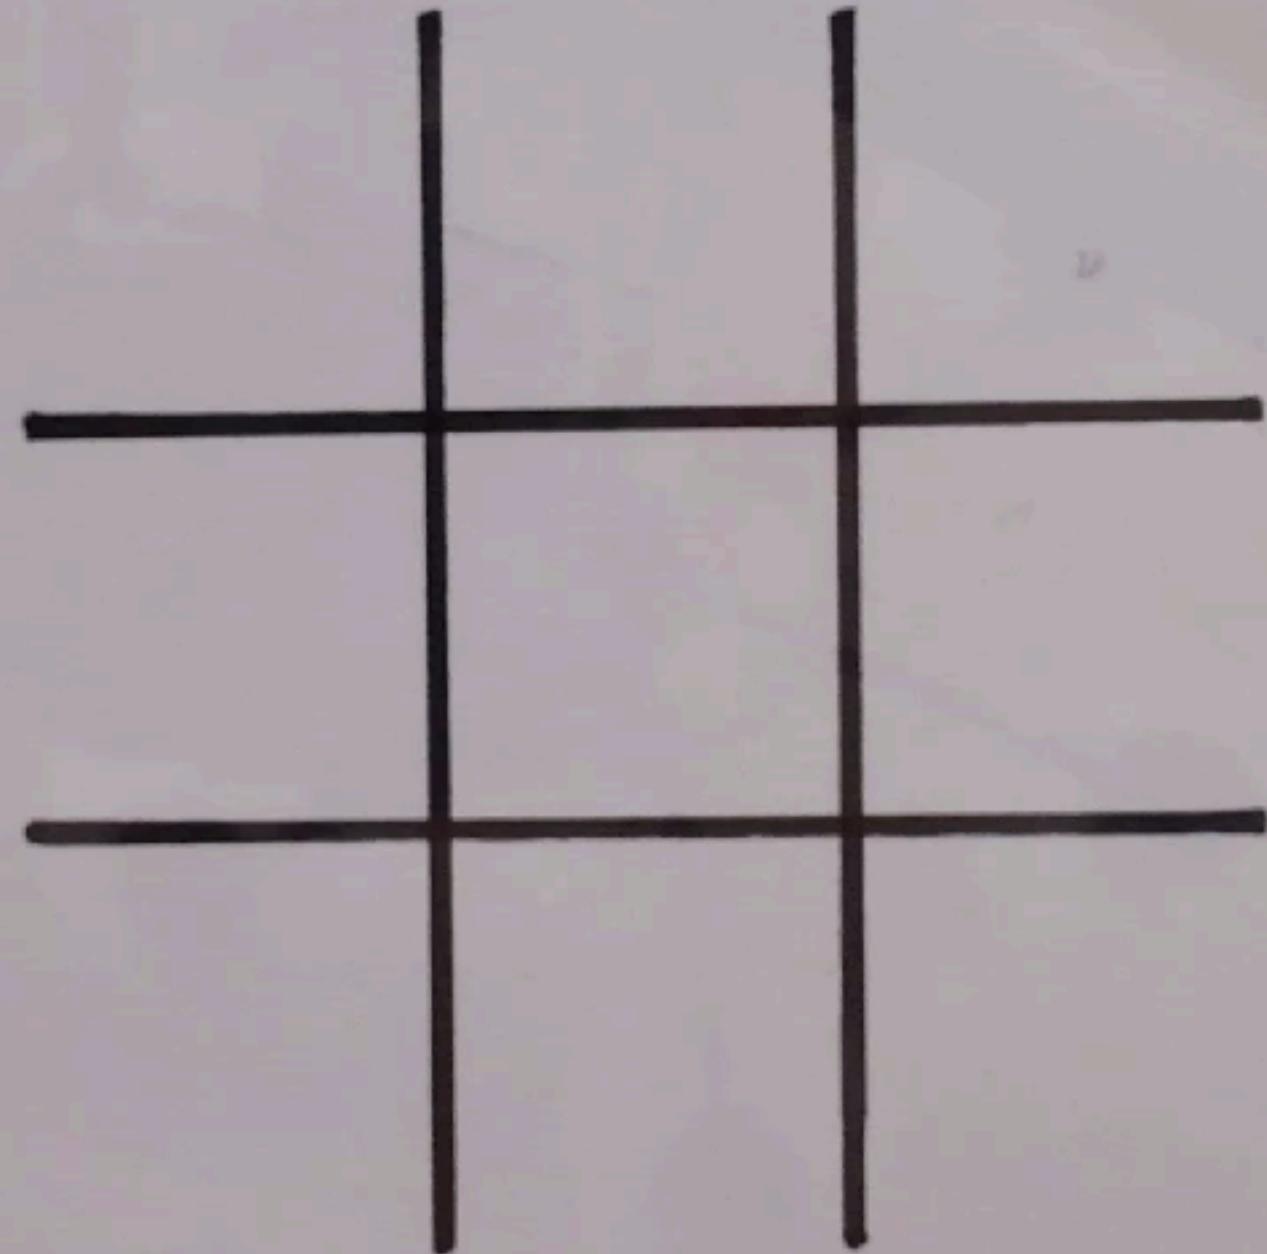
**Ephemeral** does not have a time limit.

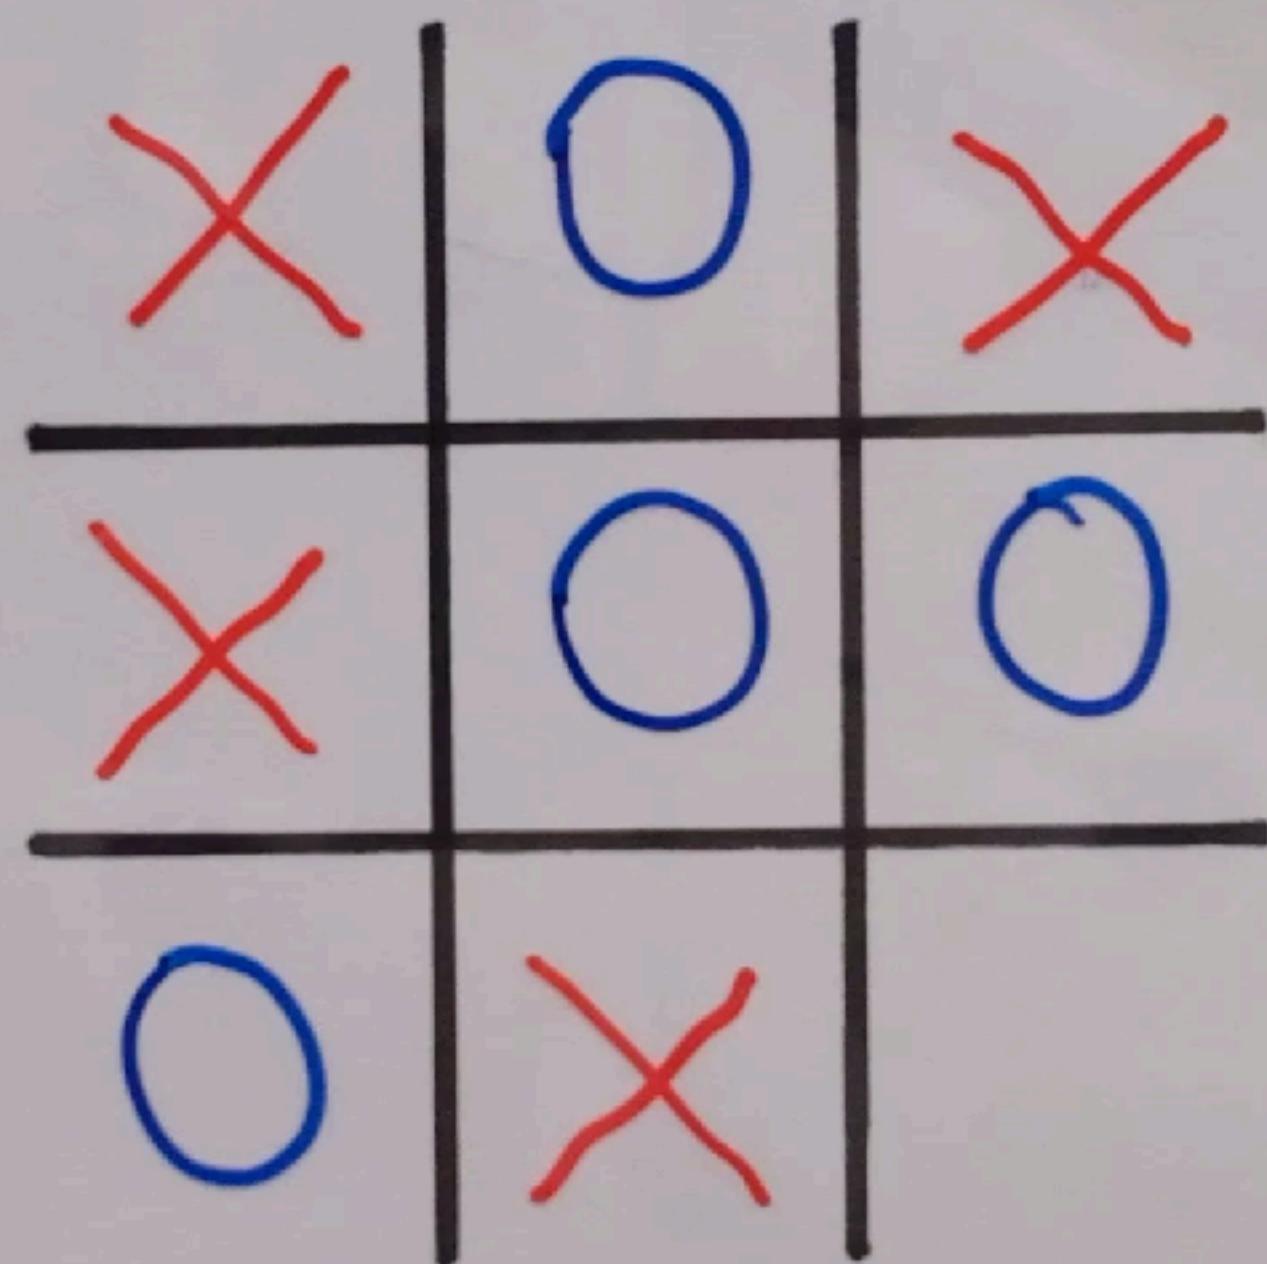
**Ephemeral** does not have a time limit.

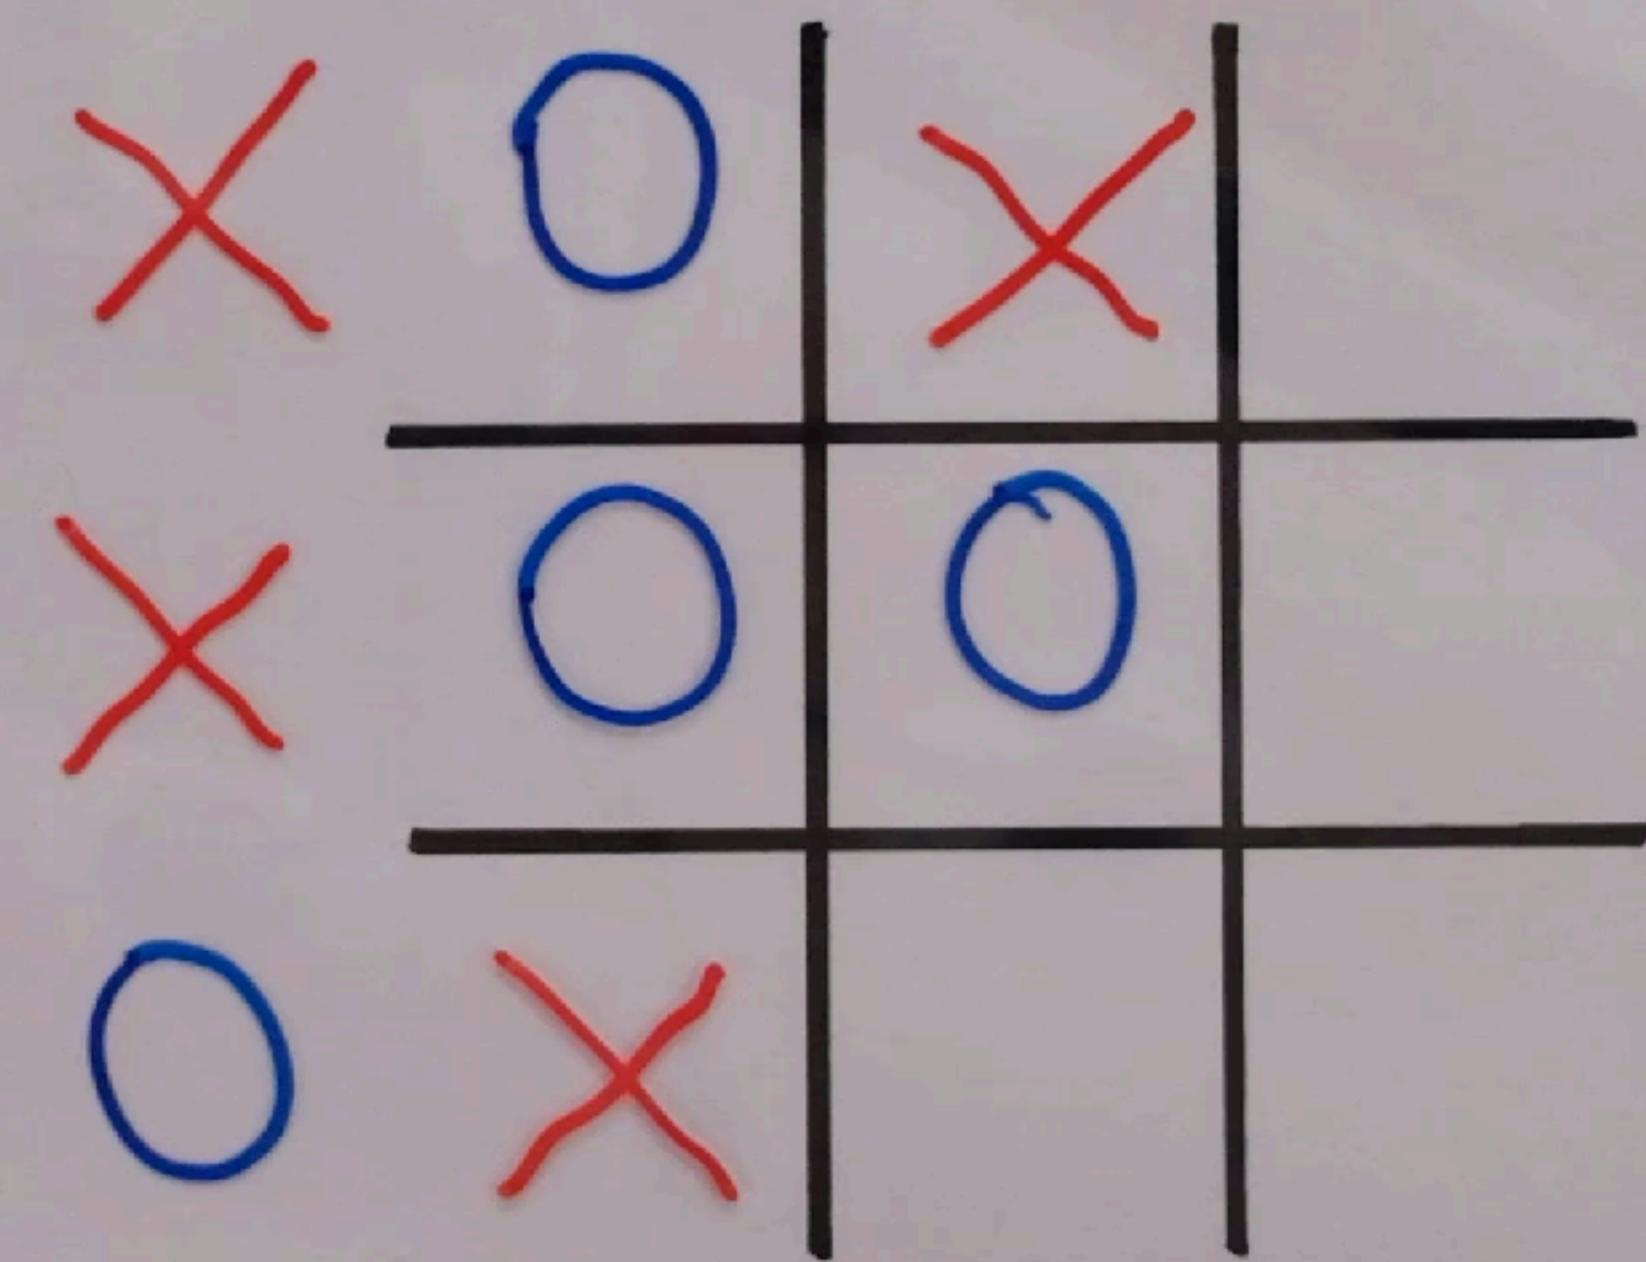
On a long enough scale,  
everything is ephemeral.

Images are **immutable** and *never change*.

# Stalemate!







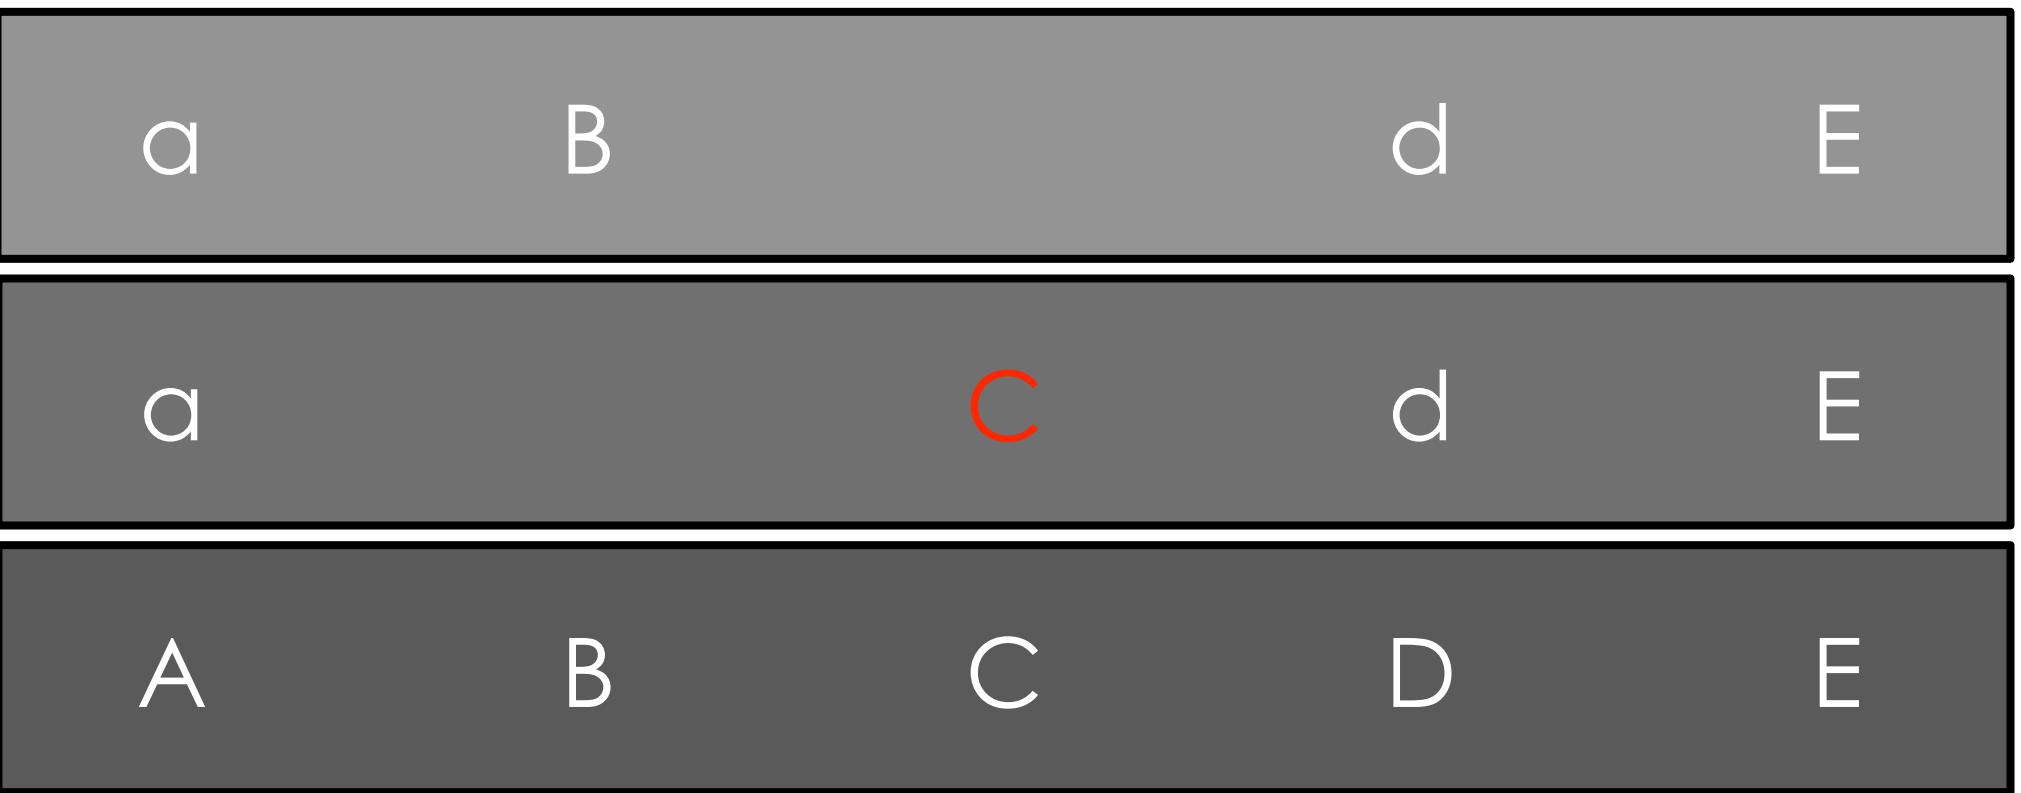
# Volumes

**Volumes** "externalize" container directories.

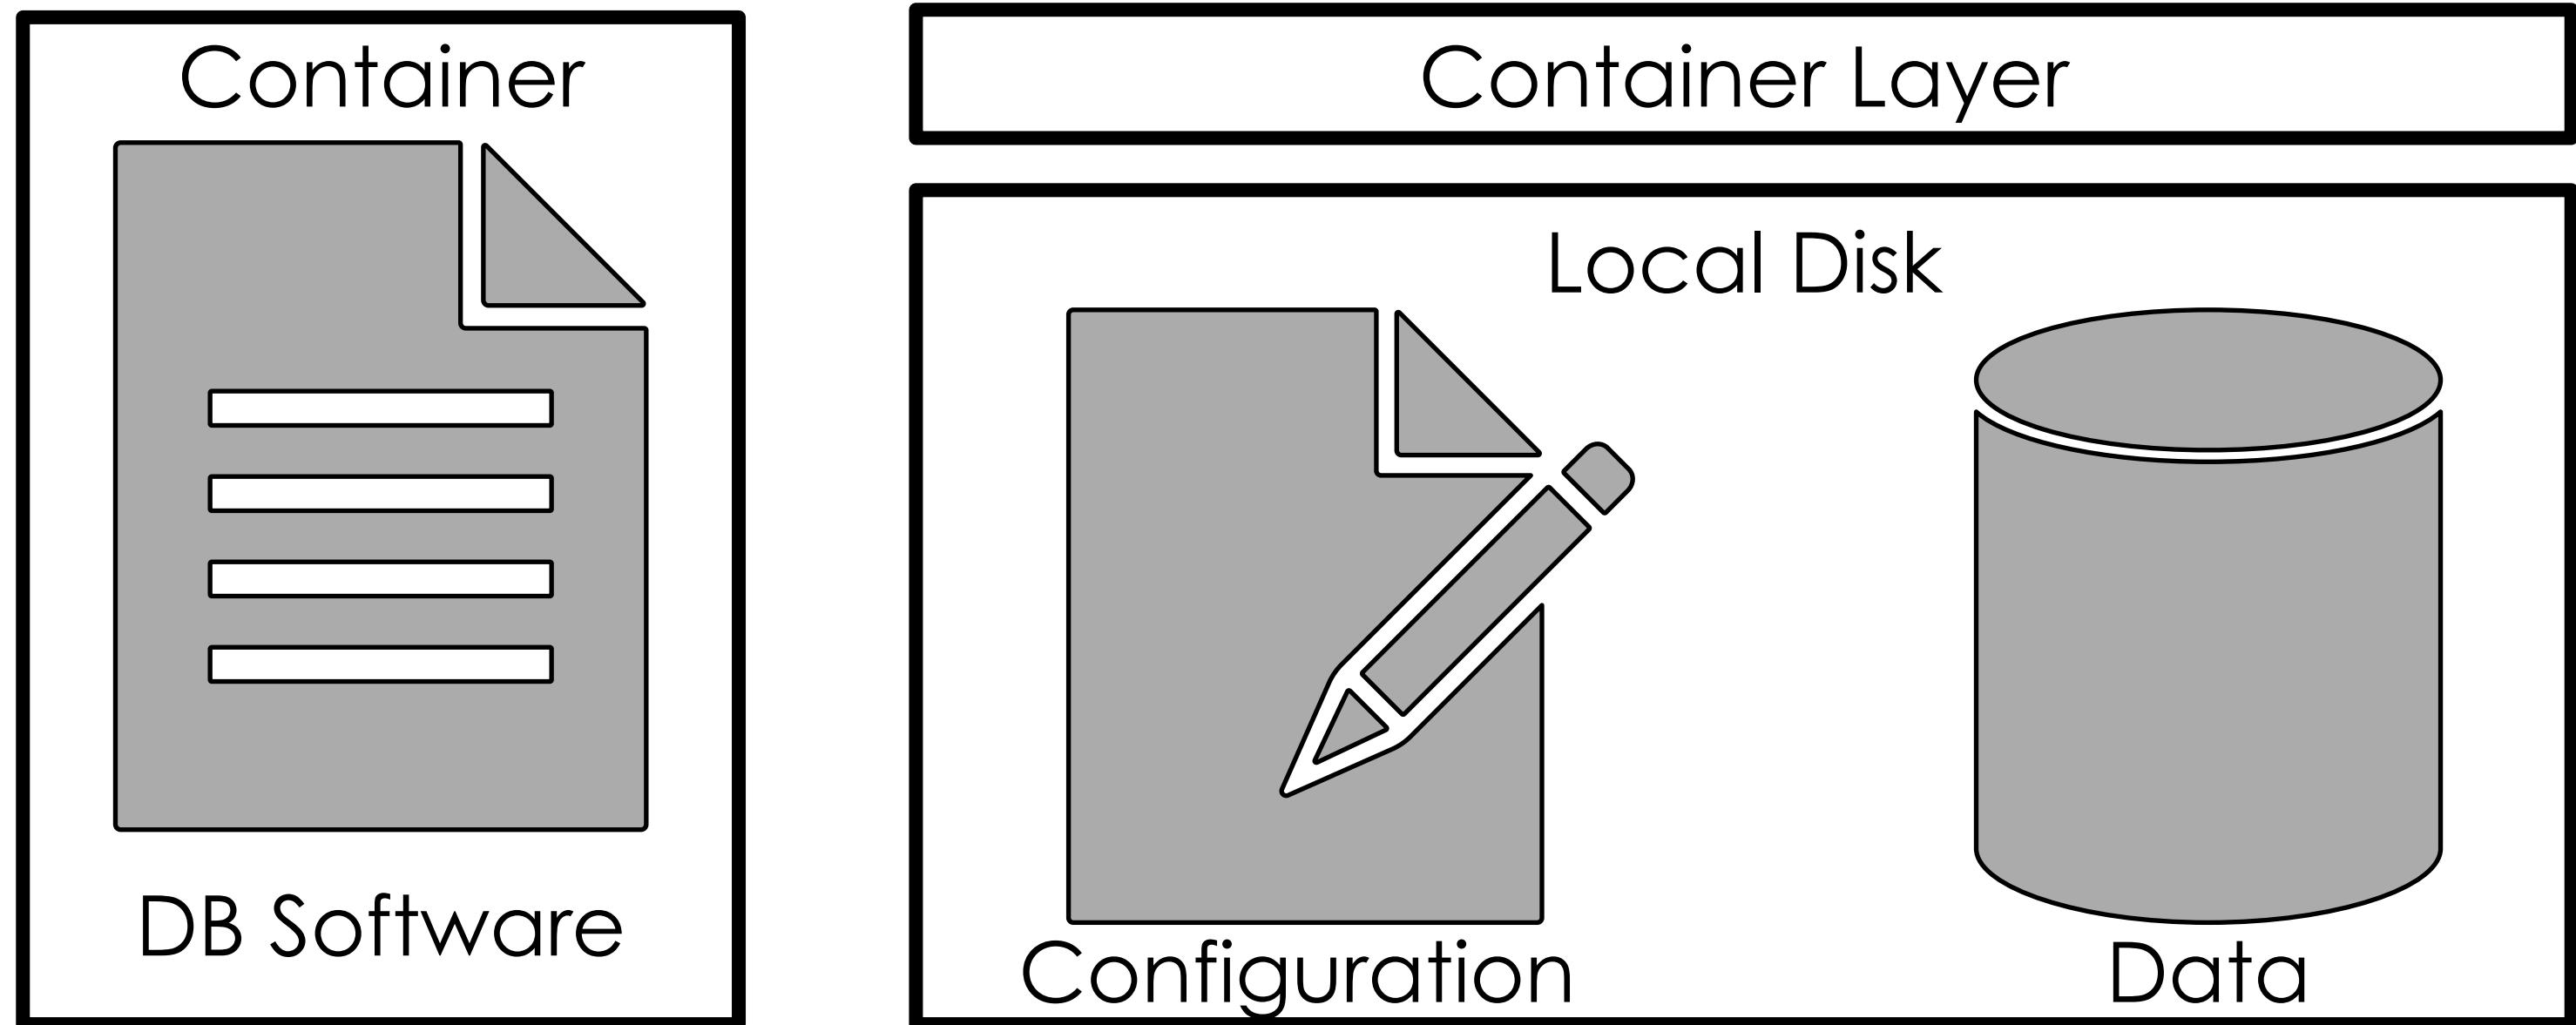
Union Layer

Container Layer

Image Layer



# Database Containers



# **Database Image Rule #1:**

If there is a database, start it.

If not, create a new database.

# Database Container Startup

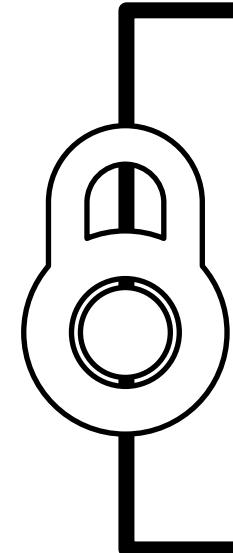
```
docker volume create ${CONTAINER_NAME}_data  
  
docker run -d \  
  --name ${CONTAINER_NAME} \  
  -v /data/${CONTAINER_NAME}:<data_directory> \  
  <image_name>
```

# Default Volumes

Linux

```
ls -l /var/lib  
drwx--x--- root root docker
```

Docker Desktop



Virtual Machine

/var/lib/docker

```
mkdir -p /data/${CONTAINER_NAME}

docker volume create \
  --opt type=none \
  --opt o=bind \
  --opt device=/data/${CONTAINER_NAME} \
${CONTAINER_NAME}_data
```

# Database Container Startup

```
docker run -d \  
  --name ${CONTAINER_NAME} \  
  -v /data/${CONTAINER_NAME}:<data_directory> \  
  <image_name>
```

# **Database Image Rule #1:**

If there is a database, start it.

If not, create a new database.

# Clone the Data Directory

```
cp -rp /data/${CONTAINER_NAME}/ \
    /data/clone
```

# Clone the Data Directory

```
cp -rp /data/clone/ \
    /data/${CONTAINER_NAME}
```

# Clone the Data Directory

```
cp -rp /data/clone/ \
    /data/${CONTAINER_NAME}
```

```
cp -rp /NFS/gold_data/clone/ \
    /data/${CONTAINER_NAME}
```

# Clone the Data Directory

```
cp -rp /data/clone/ \
    /data/${CONTAINER_NAME}
```

```
cp -rp /NFS/gold_data/clone/ \
    /data/${CONTAINER_NAME}
```

```
cp -rp /NFS/gold_data/clone/2024-05-08.0900/ \
    /data/${CONTAINER_NAME}
```

# Building Images

A **Containerfile** is a recipe for building images.

Tic-Tac-Toe Grid

Paper

**Image layers** are *modular* and *reusable*.

Layers can be **cached** and reused  
among multiple images.

Game of Sums

A Tiny Sudoku

Tic-Tac-Toe Grid

Paper

```
# Create some directories:
```

```
mkdir -p ~/workshop/  
{image,con3_upper,con3_work,container3}
```

```
# cd into the main directory:
```

```
cd ~/workshop
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=container1, \
upperdir=con3_upper, \
workdir=con3_work \
overlay container3
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=container1, \
upperdir=con3_upper, \
workdir=con3_work \
overlay container3
```

```
# Create an overlay filesystem
```

```
sudo mount -t overlay \
-o lowerdir=container1, \
upperdir=con3_upper, \
workdir=con3_work \
overlay container3
```

# What files are in each layer?

```
ls -l ~/workshop/con3_upper
```

```
ls -l ~/workshop/con3_work
```

```
ls -l ~/workshop/container3
```

```
cat ~/workshop/container3/AAA
```

```
# Update PATH
```

```
PATH=$PATH:/bin
```

```
# "Unshare" namespaces and chroot to the "container":  
  
unshare --fork --uts --pid --mount \  
    --user --ipc --net \  
    --map-root-user \  
    chroot ~/workshop/container3 \  
    /bin/sh
```

```
# Explore the container contents
```

```
ls -l
```

```
cat AAA
```

# Cleanup

```
# Cleanup
```

```
sudo umount container3
```

```
sudo umount container1
```

```
cd ~
```

```
rm -fr ~/workshop
```

# Questions

[sean.scott@viscosityna.com](mailto:sean.scott@viscosityna.com)  
<https://linktr.ee/oraclesean>

[www.viscosityna.com](http://www.viscosityna.com)



## Contact Me

[sean.scott@viscosityna.com](mailto:sean.scott@viscosityna.com)  
<https://linktr.ee/oraclesean>



